

DATASCI W261: Machine Learning at Scale

W261-1 Spring 2016

HW 12: Criteo CTR Project

April 14, 2015

Student name **Nick Hamlin**



Click-Through Rate Prediction Lab

This lab covers the steps for creating a click-through rate (CTR) prediction pipeline. You will work with the **Criteo Labs** (<http://labs.criteo.com/>) dataset that was used for a recent **Kaggle** competition (<https://www.kaggle.com/c/criteo-display-ad-challenge>).

This lab will cover:

- #####Part 1: Featurize categorical data using one-hot-encoding (OHE)
- #####Part 2: Construct an OHE dictionary
- #####Part 3: Parse CTR data and generate OHE features
 - **Visualization 1: Feature frequency**
- #####Part 4: CTR prediction and logloss evaluation
 - **Visualization 2: ROC curve**
- #####Part 5: Reduce feature dimension via feature hashing
 - **Visualization 3: Hyperparameter heat map**

Note that, for reference, you can look up the details of the relevant Spark methods in [Spark's Python API](https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD) (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD>) and the relevant NumPy methods in the [NumPy Reference](http://docs.scipy.org/doc/numpy/reference/index.html) (<http://docs.scipy.org/doc/numpy/reference/index.html>)

```
In [ ]: labVersion = 'MIDS_MLS_week12_v_0_9'
```

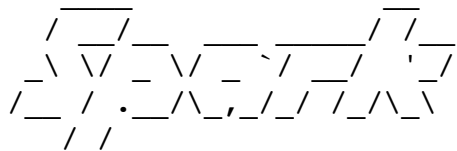
```
In [1]: #Use this to make sure we reload the MrJob code when we make changes
%load_ext autoreload
%autoreload 2
#Render matplotlib charts in notebook
%matplotlib inline

#Import some modules we know we'll use frequently
import numpy as np
import pylab as plt
```

```
In [2]: import os
import sys #current as of 9/26/2015
spark_home = os.environ['SPARK_HOME'] = \
    '/Users/nicholashamlin/spark-1.6.1-bin-hadoop2.6/'

if not spark_home:
    raise ValueError('SPARK_HOME enviroment variable is not set')
sys.path.insert(0,os.path.join(spark_home,'python'))
sys.path.insert(0,os.path.join(spark_home,'python/lib/py4j-0.9-src.zip')
execfile(os.path.join(spark_home,'python/pyspark/shell.py'))
```

Welcome to

The Spark logo is a stylized representation of the word "Spark" using a series of connected line segments that form a jagged, flame-like shape. It is positioned to the left of the text "version 1.6.1".

version 1.6.1

Using Python version 2.7.10 (default, Oct 19 2015 18:31:17)
SparkContext available as sc, HiveContext available as sqlContext.

Part 1: Featurize categorical data using one-hot-encoding

(1a) One-hot-encoding

We would like to develop code to convert categorical features to numerical ones, and to build intuition, we will work with a sample unlabeled dataset with three data points, with each data point representing an animal. The first feature indicates the type of animal (bear, cat, mouse); the second feature describes the animal's color (black, tabby); and the third (optional) feature describes what the animal eats (mouse, salmon).

In a one-hot-encoding (OHE) scheme, we want to represent each tuple of (featureID, category) via its own binary feature. We can do this in Python by creating a dictionary that maps each tuple to a distinct integer, where the integer corresponds to a binary feature. To start, manually enter the entries in the OHE dictionary associated with the sample dataset by mapping the tuples to consecutive integers starting from zero, ordering the tuples first by featureID and next by category.

Later in this lab, we'll use OHE dictionaries to transform data points into compact lists of features that can be used in machine learning algorithms.

```
In [3]: # Data for manual OHE
# Note: the first data point does not include any value for the optional
sampleOne = [(0, 'mouse'), (1, 'black')]
sampleTwo = [(0, 'cat'), (1, 'tabby'), (2, 'mouse')]
sampleThree = [(0, 'bear'), (1, 'black'), (2, 'salmon')]
sampleDataRDD = sc.parallelize([sampleOne, sampleTwo, sampleThree])
```

```
In [4]: sampleOHEDictManual = {}
sampleOHEDictManual[(0, 'bear')] = 0
sampleOHEDictManual[(0, 'cat')] = 1
sampleOHEDictManual[(0, 'mouse')] = 2
sampleOHEDictManual[(1, 'black')] = 3
sampleOHEDictManual[(1, 'tabby')] = 4
sampleOHEDictManual[(2, 'mouse')] = 5
sampleOHEDictManual[(2, 'salmon')] = 6
```

In [5]:



```

# A testing helper
#https://pypi.python.org/pypi/test_helper/0.2
import hashlib

class TestFailure(Exception):
    pass
class PrivateTestFailure(Exception):
    pass

class Test(object):
    passed = 0
    numTests = 0
    failFast = False
    private = False

    @classmethod
    def setFailFast(cls):
        cls.failFast = True

    @classmethod
    def setPrivateMode(cls):
        cls.private = True

    @classmethod
    def assertTrue(cls, result, msg=""):
        cls.numTests += 1
        if result == True:
            cls.passed += 1
            print "1 test passed."
        else:
            print "1 test failed. " + msg
            if cls.failFast:
                if cls.private:
                    raise PrivateTestFailure(msg)
                else:
                    raise TestFailure(msg)

    @classmethod
    def assertEquals(cls, var, val, msg=""):
        cls.assertTrue(var == val, msg)

    @classmethod
    def assertEqualsHashed(cls, var, hashed_val, msg=""):
        cls.assertEquals(cls._hash(var), hashed_val, msg)

    @classmethod
    def printStats(cls):
        print "{0} / {1} test(s) passed.".format(cls.passed, cls.numTests)

    @classmethod
    def _hash(cls, x):
        return hashlib.shal(str(x)).hexdigest()

```

```
In [6]: # TEST One-hot-encoding (1a)
```

```
# We have the Test class already defined above, so no need to import an
# from test_helper import Test
```

```
Test.assertEqualsHashed(sampleOHEDictManual[(0,'bear')],
                        'b6589fc6ab0dc82cf12099d1c2d40ab994e8410c',
                        "incorrect value for sampleOHEDictManual[(0,'be
Test.assertEqualsHashed(sampleOHEDictManual[(0,'cat')],
                        '356a192b7913b04c54574d18c28d46e6395428ab',
                        "incorrect value for sampleOHEDictManual[(0,'ca
Test.assertEqualsHashed(sampleOHEDictManual[(0,'mouse')],
                        'da4b9237baccdf19c0760cab7aec4a8359010b0',
                        "incorrect value for sampleOHEDictManual[(0,'mo
Test.assertEqualsHashed(sampleOHEDictManual[(1,'black')],
                        '77de68daecd823babbb58edb1c8e14d7106e83bb',
                        "incorrect value for sampleOHEDictManual[(1,'bl
Test.assertEqualsHashed(sampleOHEDictManual[(1,'tabby')],
                        '1b6453892473a467d07372d45eb05abc2031647a',
                        "incorrect value for sampleOHEDictManual[(1,'ta
Test.assertEqualsHashed(sampleOHEDictManual[(2,'mouse')],
                        'ac3478d69a3c81fa62e60f5c3696165a4e5e6ac4',
                        "incorrect value for sampleOHEDictManual[(2,'mo
Test.assertEqualsHashed(sampleOHEDictManual[(2,'salmon')],
                        'c1dfd96eea8cc2b62785275bca38ac261256e278',
                        "incorrect value for sampleOHEDictManual[(2,'sa
Test.assertEquals(len(sampleOHEDictManual.keys()), 7,
                  'incorrect number of keys in sampleOHEDictManual')
```

```
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

(1b) Sparse vectors

Data points can typically be represented with a small number of non-zero OHE features relative to the total number of features that occur in the dataset. By leveraging this sparsity and using sparse representations of OHE data, we can reduce storage and computational burdens. Below are sample vectors represented as dense numpy arrays. Use `SparseVector` (<https://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.linalg>) to represent them in a sparse fashion, and verify that both the sparse and dense representations give the same results when computing dot products (http://en.wikipedia.org/wiki/Dot_product). MLlib will need to compute dot products between `SparseVectors` and dense parameter vectors).

Use `SparseVector(size, *args)` to create a new sparse vector where `size` is the length of the vector and `args` is either a dictionary, a list of (index, value) pairs, or two separate arrays of indices and values (sorted by index). You'll need to create a sparse vector representation of each dense

aDense and bDense.

```
In [7]: import numpy as np
        from pyspark.mllib.linalg import SparseVector
```

```
In [8]: aDense = np.array([0., 3., 0., 4.])
        aSparse = SparseVector(4,[1,3],[3.,4.])

        bDense = np.array([0., 0., 0., 1.])
        bSparse = SparseVector(4,[3],[1.])

        w = np.array([0.4, 3.1, -1.4, -.5])
        print aDense.dot(w)
        print aSparse.dot(w)
        print bDense.dot(w)
        print bSparse.dot(w)

7.3
7.3
-0.5
-0.5
```

```
In [9]: # TEST Sparse Vectors (1b)
Test.assertTrue(isinstance(aSparse, SparseVector), 'aSparse needs to be
Test.assertTrue(isinstance(bSparse, SparseVector), 'aSparse needs to be
Test.assertTrue(aDense.dot(w) == aSparse.dot(w),
                'dot product of aDense and w should equal dot product o
Test.assertTrue(bDense.dot(w) == bSparse.dot(w),
                'dot product of bDense and w should equal dot product o

1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

(1c) OHE features as sparse vectors

Now let's see how we can represent the OHE features for points in our sample dataset. Using the mapping defined by the OHE dictionary from Part (1a), manually define OHE features for the three sample data points using SparseVector format. Any feature that occurs in a point should have the value 1.0. For example, the DenseVector for a point with features 2 and 4 would be [0.0, 0.0, 1.0, 0.0, 1.0, 0.0, 0.0].

```
In [ ]: # Reminder of the sample features
        # sampleOne = [(0, 'mouse'), (1, 'black')]
        # sampleTwo = [(0, 'cat'), (1, 'tabby'), (2, 'mouse')]
        # sampleThree = [(0, 'bear'), (1, 'black'), (2, 'salmon')]
```

```
In [10]: sampleOneOHEFeatManual = SparseVector(7,[2,3],[1.,1.])
sampleTwoOHEFeatManual = SparseVector(7,[1,4,5],[1.,1.,1.])
sampleThreeOHEFeatManual = SparseVector(7,[0,3,6],[1.,1.,1.])
```

```
In [11]: # TEST OHE Features as sparse vectors (1c)
Test.assertTrue(isinstance(sampleOneOHEFeatManual, SparseVector),
                  'sampleOneOHEFeatManual needs to be a SparseVector')
Test.assertTrue(isinstance(sampleTwoOHEFeatManual, SparseVector),
                  'sampleTwoOHEFeatManual needs to be a SparseVector')
Test.assertTrue(isinstance(sampleThreeOHEFeatManual, SparseVector),
                  'sampleThreeOHEFeatManual needs to be a SparseVector')
Test.assertEqualHashed(sampleOneOHEFeatManual,
                        'ecc00223d141b7bd0913d52377cee2cf5783abd6',
                        'incorrect value for sampleOneOHEFeatManual')
Test.assertEqualHashed(sampleTwoOHEFeatManual,
                        '26b023f4109e3b8ab32241938e2e9b9e9d62720a',
                        'incorrect value for sampleTwoOHEFeatManual')
Test.assertEqualHashed(sampleThreeOHEFeatManual,
                        'c04134fd603ae115395b29dcabe9d0c66fbdc8a7',
                        'incorrect value for sampleThreeOHEFeatManual')

1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

(1d) Define a OHE function

Next we will use the OHE dictionary from Part (1a) to programatically generate OHE features from the original categorical data. First write a function called `oneHotEncoding` that creates OHE feature vectors in `SparseVector` format. Then use this function to create OHE features for the first sample data point and verify that the result matches the result from Part (1c).

```
In [13]: numSampleOHEFeats = len(sampleOHEDictManual)
```



```
In [14]: def oneHotEncoding(rawFeats, OHEDict=sampleOHEDictManual, numOHEFeats=n
        """Produce a one-hot-encoding from a list of features and an OHE di

    Note:
        You should ensure that the indices used to create a SparseVecto

    Args:
        rawFeats (list of (int, str)): The features corresponding to a
            feature consists of a tuple of featureID and the feature's
        OHEDict (dict): A mapping of (featureID, value) to unique integ
        numOHEFeats (int): The total number of unique OHE features (com
            value).

    Returns:
        SparseVector: A SparseVector of length numOHEFeats with indicie
            identifiers for the (featureID, value) combinations that oc
            with values equal to 1.0.
    """
    positions=[]
    for feat in rawFeats:
        positions.append(OHEDict[feat])
    positions.sort()
    output= SparseVector(numOHEFeats,positions,[1.0]*len(positions))
    return output

# Calculate the number of features in sampleOHEDictManual
numSampleOHEFeats = len(sampleOHEDictManual)

# Run oneHotEnoding on sampleOne
sampleOneOHEFeat = oneHotEncoding(sampleOne,sampleOHEDictManual,numSamp

print sampleOneOHEFeat

(7,[2,3],[1.0,1.0])
```

```
In [15]: # TEST Define an OHE Function (1d)
Test.assertTrue(sampleOneOHEFeat == sampleOneOHEFeatManual,
                'sampleOneOHEFeat should equal sampleOneOHEFeatManual')
Test.assertEquals(sampleOneOHEFeat, SparseVector(7, [2,3], [1.0,1.0]),
                'incorrect value for sampleOneOHEFeat')
Test.assertEquals(oneHotEncoding([(1, 'black'), (0, 'mouse')], sampleOH
                numSampleOHEFeats), SparseVector(7, [2
                'incorrect definition for oneHotEncoding')

1 test passed.
1 test passed.
1 test passed.
```

(1e) Apply OHE to a dataset

Finally, use the function from Part (1d) to create OHE features for all 3 data points in the sample dataset.

```
In [16]: sampleOHEData = sampleDataRDD.map(oneHotEncoding)
        print sampleOHEData.collect()

[SparseVector(7, {2: 1.0, 3: 1.0}), SparseVector(7, {1: 1.0, 4: 1.0,
5: 1.0}), SparseVector(7, {0: 1.0, 3: 1.0, 6: 1.0})]

In [17]: # TEST Apply OHE to a dataset (1e)
        sampleOHEDataValues = sampleOHEData.collect()
        Test.assertTrue(len(sampleOHEDataValues) == 3, 'sampleOHEData should ha
        Test.assertEquals(sampleOHEDataValues[0], SparseVector(7, {2: 1.0, 3: 1
        'incorrect OHE for first sample')
        Test.assertEquals(sampleOHEDataValues[1], SparseVector(7, {1: 1.0, 4: 1
        'incorrect OHE for second sample')
        Test.assertEquals(sampleOHEDataValues[2], SparseVector(7, {0: 1.0, 3: 1
        'incorrect OHE for third sample')

1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

Part 2: Construct an OHE dictionary

(2a) Pair RDD of (featureID, category)

To start, create an RDD of distinct (featureID, category) tuples. In our sample dataset the 7 items in the resulting RDD are (0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'), (1, 'tabby'), (2, 'mouse'), (2, 'salmon'). Notably 'black' appears twice in the dataset but only contributes one item to the RDD: (1, 'black'), while 'mouse' appears twice and contributes two items: (0, 'mouse') and (2, 'mouse'). Use `flatMap` (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.flatMap>) and `distinct` (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.distinct>).

```
In [18]: def mapper(line):
        for element in line:
            yield element

        sampleDistinctFeats = (sampleDataRDD.flatMap(mapper)
                                .distinct()
                                )

        print sampleDistinctFeats
```

PythonRDD[6] at RDD at PythonRDD.scala:43

```
In [19]: # TEST Pair RDD of (featureID, category) (2a)
Test.assertEquals(sorted(sampleDistinctFeats.collect()),
                  [(0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'),
                   (1, 'tabby'), (2, 'mouse'), (2, 'salmon')],
                  'incorrect value for sampleDistinctFeats')

1 test passed.
```

(2b) OHE Dictionary from distinct features

Next, create an RDD of key-value tuples, where each `(featureID, category)` tuple in `sampleDistinctFeats` is a key and the values are distinct integers ranging from 0 to `(num keys - 1)`. Then convert this RDD into a dictionary, which can be done using the `collectAsMap` action. Note that there is no unique mapping from keys to values, as all we require is that `(featureID, category)` key be mapped to a unique integer between 0 and the number of keys. In this exercise, any valid mapping is acceptable. Use [zipWithIndex](https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.zipWithIndex) (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.zipWithIndex>) followed by `collectAsMap` (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.collectAsMap>).

In our sample dataset, one valid list of key-value tuples is: `[(0, 'bear'), 0], [(2, 'salmon'), 1], [(1, 'tabby'), 2], [(2, 'mouse'), 3], [(0, 'mouse'), 4], [(1, 'cat'), 5], [(1, 'black'), 6]`. The dictionary defined in Part (1a) illustrates another mapping between keys and integers.

```
In [20]: sampleOHEDict = (sampleDistinctFeats.zipWithIndex().collectAsMap())
print sampleOHEDict

{(2, 'mouse'): 1, (0, 'cat'): 5, (0, 'bear'): 0, (2, 'salmon'): 2,
 (1, 'tabby'): 3, (1, 'black'): 6, (0, 'mouse'): 4}
```

```
In [21]: # TEST OHE Dictionary from distinct features (2b)
Test.assertEquals(sorted(sampleOHEDict.keys()),
                  [(0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'),
                   (1, 'tabby'), (2, 'mouse'), (2, 'salmon')],
                  'sampleOHEDict has unexpected keys')
Test.assertEquals(sorted(sampleOHEDict.values()), range(7), 'sampleOHEDict has unexpected values')

1 test passed.
1 test passed.
```

(2c) Automated creation of an OHE dictionary

Now use the code from Parts (2a) and (2b) to write a function that takes an input dataset and outputs an OHE dictionary. Then use this function to create an OHE dictionary for the sample dataset, and verify that it matches the dictionary from Part (2b).

```

In [22]: def mapper(line):
          for element in line:
              yield element

def createOneHotDict(inputData):
    """Creates a one-hot-encoder dictionary based on the input data.

    Args:
        inputData (RDD of lists of (int, str)): An RDD of observations
            made up of a list of (featureID, value) tuples.

    Returns:
        dict: A dictionary where the keys are (featureID, value) tuples
            unique integers.
    """
    return inputData.flatMap(mapper).distinct().zipWithIndex().collectA

sampleOHEDictAuto = createOneHotDict(sampleDataRDD)
print sampleOHEDictAuto

{(2, 'mouse'): 1, (0, 'cat'): 5, (0, 'bear'): 0, (2, 'salmon'): 2,
(1, 'tabby'): 3, (1, 'black'): 6, (0, 'mouse'): 4}

```

```

In [23]: # TEST Automated creation of an OHE dictionary (2c)
Test.assertEquals(sorted(sampleOHEDictAuto.keys()),
                  [(0, 'bear'), (0, 'cat'), (0, 'mouse'), (1, 'black'),
                   (1, 'tabby'), (2, 'mouse'), (2, 'salmon')],
                  'sampleOHEDictAuto has unexpected keys')
Test.assertEquals(sorted(sampleOHEDictAuto.values()), range(7),
                  'sampleOHEDictAuto has unexpected values')

1 test passed.
1 test passed.

```

Part 3: Parse CTR data and generate OHE features

Before we can proceed, you'll first need to obtain the data from Criteo. If you have already completed this step in the setup lab, just run the cells below and the data will be loaded into the `rawData` variable.

Below is Criteo's data sharing agreement. After you accept the agreement, you can obtain the download URL by right-clicking on the "Download Sample" button and clicking "Copy link address" or "Copy Link Location", depending on your browser. Paste the URL into the # `TODO` cell below. The file is 8.4 MB compressed. The script below will download the file to the virtual machine (VM) and then extract the data.

If running the cell below does not render a webpage, open the [Criteo agreement](http://labs.criteo.com/downloads/2014-kaggle-display-advertising-challenge-dataset/) (<http://labs.criteo.com/downloads/2014-kaggle-display-advertising-challenge-dataset/>)

in a separate browser tab. After you accept the agreement, you can obtain the download URL by right-clicking on the "Download Sample" button and clicking "Copy link address" or "Copy Link Location", depending on your browser. Paste the URL into the # TODO cell below.

Note that the download could take a few minutes, depending upon your connection speed.

The Criteo CTR data is for HW12.1 is available here (24.3 Meg, 100,000 Rows):

https://www.dropbox.com/s/m4jlnv6rdbqzzhu/dac_sample.txt?dl=0

Alternatively you can download the sample data directly by following the instructions contained in the cell below (8M compressed).

```
In [ ]: # Run this code to view Criteo's agreement
from IPython.lib.display import IFrame

IFrame("http://labs.criteo.com/downloads/2014-kaggle-display-advertisin
        600, 350)
```

In []:

```

# Just replace <FILL IN> with the url for dac_sample.tar.gz
import glob
import os.path
import tarfile
import urllib
import urlparse

# Paste url, url should end with: dac_sample.tar.gz
url = '<FILL IN>'

url = url.strip()
baseDir = os.path.join('data')
inputPath = os.path.join('w261', 'dac_sample.txt')
fileName = os.path.join(baseDir, inputPath)
inputDir = os.path.split(fileName)[0]

def extractTar(check = False):
    # Find the zipped archive and extract the dataset
    tars = glob.glob('dac_sample*.tar.gz*')
    if check and len(tars) == 0:
        return False

    if len(tars) > 0:
        try:
            tarFile = tarfile.open(tars[0])
        except tarfile.ReadError:
            if not check:
                print 'Unable to open tar.gz file. Check your URL.'
            return False

        tarFile.extract('dac_sample.txt', path=inputDir)
        print 'Successfully extracted: dac_sample.txt'
        return True
    else:
        print 'You need to retry the download with the correct url.'
        print ('Alternatively, you can upload the dac_sample.tar.gz file to a '
              'directory')
        return False

if os.path.isfile(fileName):
    print 'File is already available. Nothing to do.'
elif extractTar(check = True):
    print 'tar.gz file was already available.'
elif not url.endswith('dac_sample.tar.gz'):
    print 'Check your download url. Are you downloading the Sample data?'
else:
    # Download the file and store it in the same directory as this note
    try:
        urllib.urlretrieve(url, os.path.basename(urlparse.urlsplit(url)))
    except IOError:
        print 'Unable to download and store: {0}'.format(url)

extractTar()

```

```

In [24]: import os.path
baseDir = os.path.join('data')
inputPath = os.path.join('w261', 'dac_sample.txt')
fileName = os.path.join(baseDir, inputPath)

#Simpler way
fileName='dac_sample.txt'
if os.path.isfile(fileName):
    rawData = (sc
                .textFile(fileName, 2)
                .map(lambda x: x.replace('\t', ','))) # work with either
    print rawData.take(1)

```

```

[u'0,1,1,5,0,1382,4,15,2,181,1,2,,2,68fd1e64,80e26c9b,fb936136,7b472
3c4,25c83c98,7e0ccccf,de7995b8,1f89b562,a73ee510,a8cd5504,b2cb9c98,3
7c9c164,2824a5f6,1adce6ef,8ba8b39a,891b62e7,e5ba7672,f54016b9,21ddcd
c9,b1252a9d,07b5194c,,3a171ecb,c5c50484,e8b83407,9727dd16']

```

(3a) Loading and splitting the data

We are now ready to start working with the actual CTR data, and our first task involves splitting it into training, validation, and test sets. Use the [randomSplit method](https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.randomSplit) (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.randomSplit>) with the specified weights and seed to create RDDs storing each of these datasets, and then [cache](https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.cache) (<https://spark.apache.org/docs/latest/api/python/pyspark.html#pyspark.RDD.cache>) each of these RDDs, as we will be accessing them multiple times in the remainder of this lab. Finally, compute the size of each dataset.

```

In [25]: weights = [.8, .1, .1]
seed = 42
# Use randomSplit with weights and seed
rawTrainData, rawValidationData, rawTestData = rawData.randomSplit(weights)
# Cache the data
rawTrainData.cache()
rawValidationData.cache()
rawTestData.cache()

nTrain = rawTrainData.count()
nVal = rawValidationData.count()
nTest = rawTestData.count()
print nTrain, nVal, nTest, nTrain + nVal + nTest
print rawData.take(1)

```

```

79911 10075 10014 100000
[u'0,1,1,5,0,1382,4,15,2,181,1,2,,2,68fd1e64,80e26c9b,fb936136,7b472
3c4,25c83c98,7e0ccccf,de7995b8,1f89b562,a73ee510,a8cd5504,b2cb9c98,3
7c9c164,2824a5f6,1adce6ef,8ba8b39a,891b62e7,e5ba7672,f54016b9,21ddcd
c9,b1252a9d,07b5194c,,3a171ecb,c5c50484,e8b83407,9727dd16']

```



```
In [26]: # TEST Loading and splitting the data (3a)
Test.assertTrue(all([rawTrainData.is_cached, rawValidationData.is_cache
                    'you must cache the split data'])
Test.assertEquals(nTrain, 79911, 'incorrect value for nTrain')
Test.assertEquals(nVal, 10075, 'incorrect value for nVal')
Test.assertEquals(nTest, 10014, 'incorrect value for nTest')

1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

(3b) Extract features

We will now parse the raw training data to create an RDD that we can subsequently use to create an OHE dictionary. Note from the `take()` command in Part (3a) that each raw data point is a string containing several fields separated by some delimiter. For now, we will ignore the first field (which is the 0-1 label), and parse the remaining fields (or raw features). To do this, complete the implementation of the `parsePoint` function.

```

In [27]: def parsePoint(point):
    """Converts a comma separated string into a list of (featureID, val

    Note:
        featureIDs should start at 0 and increase to the number of feat

    Args:
        point (str): A comma separated string where the first value is
            are features.

    Returns:
        list: A list of (featureID, value) tuples.
    """
    output=[]
    features=point.split(',')
    for i,j in enumerate(features[1:]):
        output.append((i,j))
    return output

parsedTrainFeat = rawTrainData.map(parsePoint)

numCategories = (parsedTrainFeat
    .flatMap(lambda x: x)
    .distinct()
    .map(lambda x: (x[0], 1))
    .reduceByKey(lambda x, y: x + y)
    .sortByKey()
    .collect())

print numCategories[2][1]

```

855

```

In [28]: # TEST Extract features (3b)
Test.assertEquals(numCategories[2][1], 855, 'incorrect implementation o
Test.assertEquals(numCategories[32][1], 4, 'incorrect implementation of

1 test passed.
1 test passed.

```

(3c) Create an OHE dictionary from the dataset

Note that `parsePoint` returns a data point as a list of (featureID, category) tuples, which is the same format as the sample dataset studied in Parts 1 and 2 of this lab. Using this observation, create an OHE dictionary using the function implemented in Part (2c). Note that we will assume for simplicity that all features in our CTR dataset are categorical.

```
In [29]: ctrOHEDict = createOneHotDict(parsedTrainFeat)
numCtrOHEFeats = len(ctrOHEDict.keys())
print numCtrOHEFeats
print ctrOHEDict[(0, '')]
```

```
233286
36164
```

```
In [30]: # TEST Create an OHE dictionary from the dataset (3c)
Test.assertEquals(numCtrOHEFeats, 233286, 'incorrect number of features')
Test.assertTrue((0, '') in ctrOHEDict, 'incorrect features in ctrOHEDict')
```

```
1 test passed.
1 test passed.
```

(3d) Apply OHE to the dataset

Now let's use this OHE dictionary by starting with the raw training data and creating an **R** **LabeledPoint**

(<http://spark.apache.org/docs/1.3.1/api/python/pyspark.mllib.html#pyspark.mllib.regression.LabeledPoint>) objects using OHE features. To do this, complete the implementation of the `parseOHEPoint` function. Hint: `parseOHEPoint` is an extension of the `parsePoint` function from Part (3b) and it uses the `oneHotEncoding` function from Part (1d).

```
In [31]: from pyspark.mllib.regression import LabeledPoint
```

[illegible]

```
In [33]: # TEST Apply OHE to the dataset (3d)
numNZ = sum(parsedTrainFeat.map(lambda x: len(x)).take(5))
numNZAlt = sum(OHETrainData.map(lambda lp: len(lp.features.indices)).ta
Test.assertEquals(numNZ, numNZAlt, 'incorrect implementation of parseOH
Test.assertTrue(withOneHot, 'oneHotEncoding not present in parseOHEPoin

1 test passed.
1 test passed.
```

Visualization 1: Feature frequency

We will now visualize the number of times each of the 233,286 OHE features appears in the training data. We first compute the number of times each feature appears, then bucket the features by these counts. The buckets are sized by powers of 2, so the first bucket corresponds to features that appear exactly once (2^0), the second to features that appear twice (2^1), the third to features that occur between three and four (2^2) times, the fifth bucket is five to eight (2^3) times and so on. The scatter plot below shows the logarithm of the bucket thresholds versus the logarithm of the number of features that have counts that fall in the buckets.

```
In [34]: def bucketFeatByCount(featCount):
    """Bucket the counts by powers of two."""
    for i in range(11):
        size = 2 ** i
        if featCount <= size:
            return size
    return -1

featCounts = (OHETrainData
    .flatMap(lambda lp: lp.features.indices)
    .map(lambda x: (x, 1))
    .reduceByKey(lambda x, y: x + y))
featCountsBuckets = (featCounts
    .map(lambda x: (bucketFeatByCount(x[1]), 1))
    .filter(lambda (k, v): k != -1)
    .reduceByKey(lambda x, y: x + y)
    .collect())
print featCountsBuckets

[(256, 748), (1024, 255), (2, 24076), (4, 16639), (32, 4755), (8, 11
440), (64, 2627), (128, 1476), (16, 7752), (512, 414), (1, 162813)]
```

```
In [35]: import matplotlib.pyplot as plt
```

```
x, y = zip(*featCountsBuckets)
x, y = np.log(x), np.log(y)
```


```
def preparePlot(xticks, yticks, figsize=(10.5, 6), hideLabels=False, gr
                gridWidth=1.0):
    """Template for generating the plot layout."""
    plt.close()
    fig, ax = plt.subplots(figsize=figsize, facecolor='white', edgecolor=
    ax.axes.tick_params(labelcolor='#999999', labelsizes='10')
    for axis, ticks in [(ax.get_xaxis(), xticks), (ax.get_yaxis(), ytic
        axis.set_ticks_position('none')
        axis.set_ticks(ticks)
        axis.label.set_color('#999999')
        if hideLabels: axis.set_ticklabels([])
    plt.grid(color=gridColor, linewidth=gridWidth, linestyle='-')
    map(lambda position: ax.spines[position].set_visible(False), ['bott
    return fig, ax
```

```
# generate layout and plot data
```

```
fig, ax = preparePlot(np.arange(0, 10, 1), np.arange(4, 14, 2))
ax.set_xlabel(r'$\log_e(bucketSize)$'), ax.set_ylabel(r'$\log_e(countIn
plt.scatter(x, y, s=14**2, c='#d6ebf2', edgecolors='#8cbfd0', alpha=0.7
pass
```

```
/Users/nicholashamlin/anaconda/lib/python2.7/site-packages/matplotli
b/collections.py:590: FutureWarning: elementwise comparison failed;
returning scalar instead, but in the future will perform elementwise
comparison
```

```
if self._edgecolors == str('face'):
```



We naturally would like to repeat the process from Part (3d), e.g., to compute OHE features for the validation and test datasets. However, we must be careful, as some categorical values will likely appear in new data that did not exist in the training data. To deal with this situation, update the `oneHotEncoding()` function from Part (1d) to ignore previously unseen categories, and then compute OHE features for the validation data.

```
OHEValidationData = rawValidationData.map(lambda point: parseOHEPoint(p
OHEValidationData.cache()
print OHEValidationData.take(1)
```

[illegible]

```
In [37]: # TEST Handling unseen features (3e)
numNZVal = (OHEValidationData
             .map(lambda lp: len(lp.features.indices))
             .sum())
Test.assertEqual(numNZVal, 372080, 'incorrect number of features')

1 test passed.
```

Part 4: CTR prediction and logloss evaluation

(4a) Logistic regression

We are now ready to train our first CTR classifier. A natural classifier to use in this setting models the probability of a click-through event rather than returning a binary response, and probabilistic predictions are useful. First use [LogisticRegressionWithSGD](https://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.classification.LogisticRegressionWithSGD) (<https://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.classification.LogisticRegressionWithSGD>) to train a model using `OHETrainData` with the given hyperparameter configuration. [LogisticRegressionModel](https://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.classification.LogisticRegressionModel) (<https://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.classification.LogisticRegressionModel>) Next, use the `LogisticRegressionModel.weights` and `LogisticRegressionModel.intercept` to get the model's parameters. Note that these are the names of the object's attributes and should not be confused with `model.weights` for a given model.

```
In [38]: from pyspark.mllib.classification import LogisticRegressionWithSGD

# fixed hyperparameters
numIters = 50
stepSize = 10.
regParam = 1e-6
regType = 'l2'
includeIntercept = True
```

```
In [39]: model0 = LogisticRegressionWithSGD.train(OHETrainData, iterations=numIters)
sortedWeights = sorted(model0.weights)
print sortedWeights[:5], model0.intercept

[-0.45899236853575609, -0.37973707648623956, -0.36996558266753304, -0.36934962879928263, -0.32697945415010637] 0.56455084025
```

```
In [40]: # TEST Logistic regression (4a)
Test.assertTrue(np.allclose(model0.intercept, 0.56455084025), 'incorrect intercept')
Test.assertTrue(np.allclose(sortedWeights[0:5],
                             [-0.45899236853575609, -0.37973707648623956, -0.36996558266753304, -0.36934962879928263, -0.32697945415010637]), 'incorrect weights')

1 test passed.
1 test passed.
```

(4b) Log loss

Throughout this lab, we will use log loss to evaluate the quality of models. Log loss is defined as:

$$\ell_{\log}(p, y) = \begin{cases} -\log(p) & \text{if } y = 1 \\ -\log(1 - p) & \text{if } y = 0 \end{cases}$$

where p is a probability between 0 and 1 and y is a label of either 0 or 1. Log loss is a standard evaluation criterion when predicting rare-events such as click-through rate prediction (it is also the criterion used in the [Criteo Kaggle competition \(https://www.kaggle.com/c/criteo-display-ad-challenge\)](https://www.kaggle.com/c/criteo-display-ad-challenge)). Write a function to compute log loss, and evaluate it on some sample inputs.

```
In [41]: from math import log

def computeLogLoss(p, y):
    """Calculates the value of log loss for a given probability and label

    Note:
        log(0) is undefined, so when p is 0 we need to add a small value
        and when p is 1 we need to subtract a small value (epsilon) from

    Args:
        p (float): A probability between 0 and 1.
        y (int): A label. Takes on the values 0 and 1.

    Returns:
        float: The log loss value.
    """
    epsilon = 10e-12
    if y==1:
        return -log(p+epsilon)
    elif y==0:
        return -log(1-p+epsilon)

print computeLogLoss(.5, 1)
print computeLogLoss(.5, 0)
print computeLogLoss(.99, 1)
print computeLogLoss(.99, 0)
print computeLogLoss(.01, 1)
print computeLogLoss(.01, 0)
print computeLogLoss(0, 1)
print computeLogLoss(1, 1)
print computeLogLoss(1, 0)

0.69314718054
0.69314718054
0.0100503358434
4.60517018499
4.60517018499
0.0100503358434
25.3284360229
-1.00000008274e-11
25.3284360229
```

```
In [42]: # TEST Log loss (4b)
Test.assertTrue(np.allclose([computeLogLoss(.5, 1), computeLogLoss(.01,
                                [0.69314718056, 0.0100503358535, 4.60517018
                                'computeLogLoss is not correct'])
Test.assertTrue(np.allclose([computeLogLoss(0, 1), computeLogLoss(1, 1)
                                [25.3284360229, 1.00000008275e-11, 25.32843
                                'computeLogLoss needs to bound p away from 0 and 1 by e

1 test passed.
1 test passed.
```

(4c) Baseline log loss

Next we will use the function we wrote in Part (4b) to compute the baseline log loss on the training data. A very simple yet natural baseline model is one where we always make the same prediction independent of the given datapoint, setting the predicted value equal to the fraction of training points that correspond to click-through events (i.e., where the label is one). Compute this value (which is simply the mean of the training labels), and then use it to compute the training log loss for the baseline model. The log loss for multiple observations is the mean of the individual log loss values.

```
In [43]: # Note that our dataset has a very high click-through rate by design
# In practice click-through rate can be one to two orders of magnitude
classOneFracTrain = (OHETrainData.map(lambda x: x.label).sum())/OHETrainData.n
print classOneFracTrain

logLossTrBase = OHETrainData.map(lambda x: computeLogLoss(classOneFracTrain, x.label))
print 'Baseline Train Logloss = {0:.3f}\n'.format(logLossTrBase)

0.22717773523
Baseline Train Logloss = 0.536
```

```
In [44]: # TEST Baseline log loss (4c)
Test.assertTrue(np.allclose(classOneFracTrain, 0.22717773523), 'incorrect value')
Test.assertTrue(np.allclose(logLossTrBase, 0.535844), 'incorrect value')

1 test passed.
1 test passed.
```

(4d) Predicted probability

In order to compute the log loss for the model we trained in Part (4a), we need to write code to generate predictions from this model. Write a function that computes the raw linear prediction from this logistic regression model and then passes it through a sigmoid function ([http://en.wikipedia.org/wiki/Sigmoid function](http://en.wikipedia.org/wiki/Sigmoid_function)) $\sigma(t) = (1 + e^{-t})^{-1}$ to return the model's probabilistic prediction. Then compute probabilistic predictions on the training data.

Note that when incorporating an intercept into our predictions, we simply add the

intercept to the value of the prediction obtained from the weights and features. Alternatively, if the intercept was included as the first weight, we would need to add a corresponding feature to our data where the feature has the value one. This is not the case here.

```
In [106]: from math import exp #  $\exp(-t) = e^{-t}$ 

def getP(x, w, intercept):
    """Calculate the probability for an observation given a set of weights

    Note:
        We'll bound our raw prediction between 20 and -20 for numerical stability

    Args:
        x (SparseVector): A vector with values of 1.0 for features that are present in the
            observation and 0.0 otherwise.
        w (DenseVector): A vector of weights (betas) for the model.
        intercept (float): The model's intercept.

    Returns:
        float: A probability between 0 and 1.
    """
    rawPrediction = x.dot(w) + intercept

    # Bound the raw prediction value
    rawPrediction = min(rawPrediction, 20)
    rawPrediction = max(rawPrediction, -20)

    output = (1 + exp(-rawPrediction))**-1
    return output

trainingPredictions = OHETrainData.map(lambda x: getP(x.features, model0.weights))

print trainingPredictions.take(5)

[0.3026288202391113, 0.10362661997434088, 0.283634247838756, 0.17846102057880123, 0.5389775379218853]
```

```
In [107]: # TEST Predicted probability (4d)
Test.assertTrue(np.allclose(trainingPredictions.sum(), 18135.4834348),
                'incorrect value for trainingPredictions')

1 test passed.
```

(4e) Evaluate the model

We are now ready to evaluate the quality of the model we trained in Part (4a). To do this, first write a general function that takes as input a model and data, and outputs the log loss. Then run this function on the OHE training data, and compare the result with the baseline log loss.

```
In [113]: def evaluateResults(model, data):
    """Calculates the log loss for the data given the model.

    Args:
        model (LogisticRegressionModel): A trained logistic regression
        data (RDD of LabeledPoint): Labels and features for each observ

    Returns:
        float: Log loss for the data.
    """
    output=data.map(lambda x: computeLogLoss(getP(x.features,model.weig
    return output

logLossTrLR0 = evaluateResults(model0, OHETrainData)
print ('OHE Features Train Logloss:\n\tBaseline = {0:.3f}\n\tLogReg = {
    .format(logLossTrBase, logLossTrLR0))

OHE Features Train Logloss:
    Baseline = 0.536
    LogReg = 0.457
```

```
In [114]: # TEST Evaluate the model (4e)
Test.assertTrue(np.allclose(logLossTrLR0, 0.456903), 'incorrect value f
1 test passed.
```

(4f) Validation log loss

Next, following the same logic as in Parts (4c) and 4(e), compute the validation log loss for both the baseline and logistic regression models. Notably, the baseline model for the validation data should still be based on the label fraction from the training dataset.

```
In [56]: logLossValBase = OHEValidationData.map(lambda x: computeLogLoss(classOn

logLossValLR0 = evaluateResults(model0, OHEValidationData)
print ('OHE Features Validation Logloss:\n\tBaseline = {0:.3f}\n\tLogRe
    .format(logLossValBase, logLossValLR0))

OHE Features Validation Logloss:
    Baseline = 0.528
    LogReg = 0.457
```

```
In [57]: # TEST Validation log loss (4f)
Test.assertTrue(np.allclose(logLossValBase, 0.527603), 'incorrect value
Test.assertTrue(np.allclose(logLossValLR0, 0.456957), 'incorrect value

1 test passed.
1 test passed.
```

Visualization 2: ROC curve


We will now visualize how well the model predicts our target. To do this we generate a plot of the ROC curve. The ROC curve shows us the trade-off between the false positive rate and true positive rate, as we liberalize the threshold required to predict a positive outcome. A random model is represented by the dashed line.

```
In [58]: labelsAndScores = OHEValidationData.map(lambda lp:
                                                (lp.label, getP(lp.features
labelsAndWeights = labelsAndScores.collect()
labelsAndWeights.sort(key=lambda (k, v): v, reverse=True)
labelsByWeight = np.array([k for (k, v) in labelsAndWeights])

length = labelsByWeight.size
truePositives = labelsByWeight.cumsum()
numPositive = truePositives[-1]
falsePositives = np.arange(1.0, length + 1, 1.) - truePositives

truePositiveRate = truePositives / numPositive
falsePositiveRate = falsePositives / (length - numPositive)

# Generate layout and plot data
fig, ax = preparePlot(np.arange(0., 1.1, 0.1), np.arange(0., 1.1, 0.1))
ax.set_xlim(-.05, 1.05), ax.set_ylim(-.05, 1.05)
ax.set_ylabel('True Positive Rate (Sensitivity)')
ax.set_xlabel('False Positive Rate (1 - Specificity)')
plt.plot(falsePositiveRate, truePositiveRate, color='#8cbfd0', linestyle)
plt.plot((0., 1.), (0., 1.), linestyle='--', color='#d6ebf2', linewidth
pass
```



Part 5: Reduce feature dimension via feature hashing

(5a) Hash function

As we just saw, using a one-hot-encoding featurization can yield a model with good statistical accuracy. However, the number of distinct categories across all features is quite large -- recall that we observed 233K categories in the training data in Part (3c). Moreover, the full Kaggle training dataset includes more than 33M distinct categories, and the Kaggle dataset itself is just a small subset of Criteo's labeled data. Hence, featurizing via a one-hot-encoding representation would lead to a very large feature vector. To reduce the dimensionality of the feature space, we will use feature hashing.

####Below is the hash function that we will use for this part of the lab. We will first use this hash function with the three sample data points from Part (1a) to gain some intuition. Specifically, run code to hash the three sample points using two different values for numBuckets and observe the resulting hashed feature dictionaries.

```
In [59]: from collections import defaultdict
import hashlib

def hashFunction(numBuckets, rawFeats, printMapping=False):
    """Calculate a feature dictionary for an observation's features bas

    Note:
        Use printMapping=True for debug purposes and to better understand

    Args:
        numBuckets (int): Number of buckets to use as features.
        rawFeats (list of (int, str)): A list of features for an observation
            (featureID, value) tuples.
        printMapping (bool, optional): If true, the mappings of featureID to
            bucket are printed.

    Returns:
        dict of int to float: The keys will be integers which represent the
            features have been hashed to. The value for a given key will be the
            sum of (featureID, value) tuples that have hashed to that key.

    """
    mapping = {}
    for ind, category in rawFeats:
        featureString = category + str(ind)
        mapping[featureString] = int(int(hashlib.md5(featureString).hexdigest()) % numBuckets)
    if(printMapping): print mapping
    sparseFeatures = defaultdict(float)
    for bucket in mapping.values():
        sparseFeatures[bucket] += 1.0
    return dict(sparseFeatures)

# Reminder of the sample values:
# sampleOne = [(0, 'mouse'), (1, 'black')]
# sampleTwo = [(0, 'cat'), (1, 'tabby'), (2, 'mouse')]
# sampleThree = [(0, 'bear'), (1, 'black'), (2, 'salmon')]
```

```

In [60]: # Use four buckets
sampOneFourBuckets = hashFunction(4, sampleOne, True)
sampTwoFourBuckets = hashFunction(4, sampleTwo, True)
sampThreeFourBuckets = hashFunction(4, sampleThree, True)

# Use one hundred buckets
sampOneHundredBuckets = hashFunction(100, sampleOne, True)
sampTwoHundredBuckets = hashFunction(100, sampleTwo, True)
sampThreeHundredBuckets = hashFunction(100, sampleThree, True)

print '\t\t 4 Buckets \t\t\t 100 Buckets'
print 'SampleOne:\t {0}\t\t {1}'.format(sampOneFourBuckets, sampOneHund
print 'SampleTwo:\t {0}\t\t {1}'.format(sampTwoFourBuckets, sampTwoHund
print 'SampleThree:\t {0}\t {1}'.format(sampThreeFourBuckets, sampThree

{'black1': 2, 'mouse0': 3}
{'cat0': 0, 'tabby1': 0, 'mouse2': 2}
{'bear0': 0, 'black1': 2, 'salmon2': 1}
{'black1': 14, 'mouse0': 31}
{'cat0': 40, 'tabby1': 16, 'mouse2': 62}
{'bear0': 72, 'black1': 14, 'salmon2': 5}

          4 Buckets                                100 Buckets
SampleOne:      {2: 1.0, 3: 1.0}                    {14: 1.0, 31: 1.0}
SampleTwo:      {0: 2.0, 2: 1.0}                    {40: 1.0, 16: 1.0,
62: 1.0}
SampleThree:    {0: 1.0, 1: 1.0, 2: 1.0}            {72: 1.0, 5: 1.0, 1
4: 1.0}

In [61]: # TEST Hash function (5a)
Test.assertEquals(sampOneFourBuckets, {2: 1.0, 3: 1.0}, 'incorrect valu
Test.assertEquals(sampThreeHundredBuckets, {72: 1.0, 5: 1.0, 14: 1.0},
                  'incorrect value for sampThreeHundredBuckets')

1 test passed.
1 test passed.

```

(5b) Creating hashed features

Next we will use this hash function to create hashed features for our CTR datasets. First write a function that uses the hash function from Part (5a) with `numBuckets = $2^{15} \approx 33K$` to create a `LabeledPoint` with hashed features stored as a `SparseVector`. Then use this function to create new training, validation and test datasets with hashed features. Hint: `parsedHashPoint` is similar to `parseOHEPoint` from Part (3d).


```
In [81]: # TEST Creating hashed features (5b)
hashTrainDataFeatureSum = sum(hashTrainData
                                .map(lambda lp: len(lp.features.indices))
                                .take(20))
hashTrainDataLabelSum = sum(hashTrainData
                              .map(lambda lp: lp.label)
                              .take(100))
hashValidationDataFeatureSum = sum(hashValidationData
                                    .map(lambda lp: len(lp.features.indices)
                                    .take(20))
hashValidationDataLabelSum = sum(hashValidationData
                                  .map(lambda lp: lp.label)
                                  .take(100))
hashTestDataFeatureSum = sum(hashTestData
                              .map(lambda lp: len(lp.features.indices))
                              .take(20))
hashTestDataLabelSum = sum(hashTestData
                            .map(lambda lp: lp.label)
                            .take(100))

Test.assertEquals(hashTrainDataFeatureSum, 772, 'incorrect number of fe
Test.assertEquals(hashTrainDataLabelSum, 24.0, 'incorrect labels in has
Test.assertEquals(hashValidationDataFeatureSum, 776,
                  'incorrect number of features in hashValidationData')
Test.assertEquals(hashValidationDataLabelSum, 16.0, 'incorrect labels i
Test.assertEquals(hashTestDataFeatureSum, 774, 'incorrect number of fea
Test.assertEquals(hashTestDataLabelSum, 23.0, 'incorrect labels in hash

1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
1 test passed.
```

(5c) Sparsity

Since we have 33K hashed features versus 233K OHE features, we should expect OHE features to be sparser. Verify this hypothesis by computing the average sparsity of the OHE and the hashed training datasets.

Note that if you have a `SparseVector` named `sparse`, calling `len(sparse)` returns the total number of features, not the number features with entries. `SparseVector` objects have the attributes `indices` and `values` that contain information about which features are nonzero. Continuing with our example, these can be accessed using `sparse.indices` and `sparse.values`, respectively.

```
In [97]: from __future__ import division

def computeSparsity(data, d, n):
    """Calculates the average sparsity for the features in an RDD of La

    Args:
        data (RDD of LabeledPoint): The LabeledPoints to use in the spa
        d (int): The total number of features.
        n (int): The number of observations in the RDD.

    Returns:
        float: The average of the ratio of features in a point to total
    """
    return data.map(lambda x: len(x.features.indices)/d).sum()/n

averageSparsityHash = computeSparsity(hashTrainData, numBucketsCTR, nTr
averageSparsityOHE = computeSparsity(OHETrainData, numCtrOHEFeats, nTra

print 'Average OHE Sparsity: {0:.7e}'.format(averageSparsityOHE)
print 'Average Hash Sparsity: {0:.7e}'.format(averageSparsityHash)

Average OHE Sparsity: 1.6717677e-04
Average Hash Sparsity: 1.1805561e-03
```

```
In [98]: # TEST Sparsity (5c)
Test.assertTrue(np.allclose(averageSparsityOHE, 1.6717677e-04),
                'incorrect value for averageSparsityOHE')
Test.assertTrue(np.allclose(averageSparsityHash, 1.1805561e-03),
                'incorrect value for averageSparsityHash')

1 test passed.
1 test passed.
```

(5d) Logistic model with hashed features

Now let's train a logistic regression model using the hashed features. Run a grid search to find suitable hyperparameters for the hashed features, evaluating via log loss on the validation data. Note: This may take a few minutes to run. Use 1 and 10 for stepSizes and $1e-6$ and $1e-3$ for regParams.

```
In [99]: numIters = 500
         regType = 'l2'
         includeIntercept = True

         # Initialize variables using values from initial model training
         bestModel = None
         bestLogLoss = 1e10
```

```

In [115]: stepSizes = [1,10]
regParams = [1e-6,1e-3]
for stepSize in stepSizes:
    for regParam in regParams:
        model = (LogisticRegressionWithSGD
                  .train(hashTrainData, numIters, stepSize, regParam=reg
                        intercept=includeIntercept))
        logLossVa = evaluateResults(model, hashValidationData) #####
        print ('\tstepSize = {0:.1f}, regParam = {1:.0e}: logloss = {2:
              .format(stepSize, regParam, logLossVa))
        if (logLossVa < bestLogLoss):
            bestModel = model
            bestLogLoss = logLossVa

print ('Hashed Features Validation Logloss:\n\tBaseline = {0:.3f}\n\tLo
      .format(logLossValBase, bestLogLoss))

      stepSize = 1.0, regParam = 1e-06: logloss = 0.475
      stepSize = 1.0, regParam = 1e-03: logloss = 0.475
      stepSize = 10.0, regParam = 1e-06: logloss = 0.450
      stepSize = 10.0, regParam = 1e-03: logloss = 0.452
Hashed Features Validation Logloss:
      Baseline = 0.528
      LogReg = 0.450

```

```

In [118]: # TEST Logistic model with hashed features (5d)
# This unit test appears to have a slightly different log-loss value, s
Test.assertTrue(np.allclose(bestLogLoss, 0.449740139932), 'incorrect va

1 test passed.

```

Visualization 3: Hyperparameter heat map

We will now perform a visualization of an extensive hyperparameter search. Specifically, we will create a heat map where the brighter colors correspond to lower values of `logLoss`.

The search was run using six step sizes and six values for regularization, which required the training of thirty-six separate models. We have included the results below, but omitted the actual search to save time.


```
In [102]: from matplotlib.colors import LinearSegmentedColormap
```

```
# Saved parameters and results. Eliminate the time required to run 36
stepSizes = [3, 6, 9, 12, 15, 18]
regParams = [1e-7, 1e-6, 1e-5, 1e-4, 1e-3, 1e-2]
logLoss = np.array([[ 0.45808431,  0.45808493,  0.45809113,  0.45815333
                      [ 0.45188196,  0.45188306,  0.4518941,   0.4520051,
                      [ 0.44886478,  0.44886613,  0.44887974,  0.44902096
                      [ 0.44706645,  0.4470698,   0.44708102,  0.44724251
                      [ 0.44588848,  0.44589365,  0.44590568,  0.44606631
                      [ 0.44508948,  0.44509474,  0.44510274,  0.44525007

numRows, numCols = len(stepSizes), len(regParams)
logLoss = np.array(logLoss)
logLoss.shape = (numRows, numCols)

fig, ax = preparePlot(np.arange(0, numCols, 1), np.arange(0, numRows, 1
                      hideLabels=True, gridWidth=0.)
ax.set_xticklabels(regParams), ax.set_yticklabels(stepSizes)
ax.set_xlabel('Regularization Parameter'), ax.set_ylabel('Step Size')

colors = LinearSegmentedColormap.from_list('blue', ['#0022ff', '#000055']
image = plt.imshow(logLoss, interpolation='nearest', aspect='auto',
                    cmap = colors)
pass
```



(5e) Evaluate on the test set

Finally, evaluate the best model from Part (5d) on the test set. Compare the resulting log loss with the baseline log loss on the test set, which can be computed in the same way that the validation log loss was computed in Part (4f).

```
In [119]: # Log loss for the best model from (5d)
logLossTest = evaluateResults(bestModel, hashTestData)

# Log loss for the baseline model
logLossTestBaseline = hashTestData.map(lambda x: computeLogLoss(classOn

print ('Hashed Features Test Log Loss:\n\tBaseline = {0:.3f}\n\tLogReg
      .format(logLossTestBaseline, logLossTest))

Hashed Features Test Log Loss:
      Baseline = 0.537
      LogReg = 0.457

In [122]: # TEST Evaluate on the test set (5e)
Test.assertTrue(np.allclose(logLossTestBaseline, 0.537438),
                'incorrect value for logLossTestBaseline')

# This unit-test is also affected by the rounding error, so I've modifi
Test.assertTrue(np.allclose(logLossTest, 0.457255168698), 'incorrect va

1 test passed.
1 test passed.
```